

BananaNavigationSolution

April 9, 2020

```
In [1]: import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim

import numpy as np
import random
from collections import namedtuple, deque, OrderedDict
import pandas as pd
```

```
In [2]: !pip -q install ./python
```

```
tensorflow 1.7.1 has requirement numpy>=1.13.3, but you'll have numpy 1.12.1 which is incompatible
ipython 6.5.0 has requirement prompt-toolkit<2.0.0,>=1.0.15, but you'll have prompt-toolkit 3.0.0
```

```
In [3]: from unityagents import UnityEnvironment
env = UnityEnvironment(file_name="/data/Banana_Linux_NoVis/Banana.x86_64")
```

INFO:unityagents:

'Academy' started successfully!

Unity Academy name: Academy

Number of Brains: 1

Number of External Brains : 1

Lesson number : 0

Reset Parameters :

Unity brain name: BananaBrain

Number of Visual Observations (per agent): 0

Vector Observation space type: continuous

Vector Observation space size (per agent): 37

Number of stacked Vector Observation: 1

Vector Action space type: discrete

Vector Action space size (per agent): 4

Vector Action descriptions: , , ,

```
In [4]: import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [5]: # get the default brain
        brain_name = env.brain_names[0]
        brain = env.brains[brain_name]

        env_info = env.reset(train_mode=True)[brain_name]
        state_size = len(env_info.vector_observations[0])
        action_size = brain.vector_action_space_size

        print(brain_name)
```

BananaBrain

```
In [6]: def QNet(state_size=state_size, action_size=action_size, seed=10, fc1=256, fc2=512):
        torch.manual_seed(seed)
        return nn.Sequential(OrderedDict([
            ('fc1', nn.Linear(state_size, fc1)),
            ('relu1', nn.ReLU()),
            ('dropout1', nn.Dropout(p=0.25)),
            ('fc2', nn.Linear(fc1, fc2)),
            ('relu2', nn.ReLU()),
            ('dropout2', nn.Dropout(p=0.25)),

            ('output', nn.Linear(fc2, action_size))
        ]))
```

```
In [7]: BUFFER_SIZE = 100000
        BATCH_SIZE = 64
        GAMMA = 0.995
        TAU = 1e-3
        LR = 0.001
        UPDATE_EVERY = 4 # how often to update the network
```

```
In [8]: episodes = 1000
        horizon_t = 1000
```

```
In [9]: device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
        print(device)
```

cpu

```
In [10]: class ReplayBuffer():
        """Create the replay buffer to store the experience from the env as (s,a,r,s',done)

        def __init__(self, buffer_size, batch_size, seed=231):
            """Initialize ReplayBuffer.

            Params
```

```

=====
    buffer_size (int): (maximal) capacity of the buffer and size of sampled exp
    batch_size (int): size of the random samples
    seed (int): randomize
    """
    self.batch_size = batch_size
    self.memory = deque(maxlen=buffer_size)
    self.experience = namedtuple("Experience", field_names=["state", "action", "reward", "next_state", "done"])
    random.seed(seed)

def add(self, state, action, reward, next_state, done):
    """Add the tuple s,a,r,s',done in the experience replay buffer"""
    e = self.experience(state, action, reward, next_state, done)
    self.memory.append(e)

def sample(self):
    """Sample random experiences with the size of batch_size.
    Returns a tuple (s,a,r,s',done), each item as a torch vector"""

    samples = random.sample(self.memory, self.batch_size)

    #return torch tensors
    states = torch.from_numpy(np.vstack([e.state for e in samples if e is not None]))
    actions = torch.from_numpy(np.vstack([e.action for e in samples if e is not None]))
    rewards = torch.from_numpy(np.vstack([e.reward for e in samples if e is not None]))
    next_states = torch.from_numpy(np.vstack([e.next_state for e in samples if e is not None]))
    dones = torch.from_numpy(np.vstack([e.done for e in samples if e is not None])).float()
    return (states, actions, rewards, next_states, dones)

def __len__(self):
    return len(self.memory)

```

In [11]: class Agent():

```

def __init__(self, state_size, action_size, seed=231):
    self.state_size = state_size
    self.action_size = action_size
    self.qnet_local = QNet(state_size, action_size).to(device)
    self.qnet_target = QNet(state_size, action_size).to(device)
    self.replay = ReplayBuffer(buffer_size=BUFFER_SIZE, batch_size=BATCH_SIZE)
    self.t_steps = 0
    self.optimizer = optim.Adam(params=self.qnet_local.parameters(), lr=LR)
    random.seed(seed)

def act(self, state, eps):

```

```

#select the next state based on policy
state = torch.from_numpy(state).float().unsqueeze(0).to(device)
self.qnet_local.eval()
with torch.no_grad():
    action_values = self.qnet_local(state)
self.qnet_local.train()

if random.random() > eps:
    # go greedy
    #when run in gpu mode, tensor's data cannot be converted to numpy
    return np.argmax(action_values.data.cpu().numpy())
else:
    #random action
    return np.random.choice(np.arange(self.action_size))

def step(self, state, action, reward, next_state, done):
    self.replay.add(state, action, reward, next_state, done)
    self.t_steps += 1
    if self.t_steps % UPDATE_EVERY == 0:
        # check if enough samples
        if self.replay.__len__() > BATCH_SIZE:
            subset = self.replay.sample()
            self.learn(subset, GAMMA)

def learn(self, experience, gamma):
    state, action, reward, next_state, done = experience
    # Get max predicted Q values (for next states) from target model
    Q_targets_next = self.qnet_target(next_state).detach().max(1)[0].unsqueeze(1) #
    # Compute Q targets for current states
    Q_targets = reward + (gamma * Q_targets_next * (1 - done))
    Q_expected = self.qnet_local(state).gather(1, action) #view in details

    loss = nn.functional.mse_loss(Q_expected, Q_targets)

    self.optimizer.zero_grad()
    loss.backward()
    self.optimizer.step()

    self.soft_update(self.qnet_local, self.qnet_target)

def soft_update(self, local_model, target_model, tau=TAU):
    for target_param, local_param in zip(target_model.parameters(), local_model.par
        target_param.data.copy_(tau*local_param.data + (1.0-tau)*target_param.data)

```

```

In [12]: agent = Agent(state_size=state_size, action_size=action_size, seed=10)

In [13]: def dqn(episodes, horizon_t, eps_start = 0.99, eps_end = 0.01, eps_decay = 0.995):
    scores = []
    scores_win = deque(maxlen=100) #last consecutive episode scores
    eps = eps_start
    for e in range(1, episodes):
        env_info = env.reset(train_mode=True)[brain_name]
        state = env_info.vector_observations[0]
        score = 0
        for t in range(horizon_t):
            action = agent.act(state, eps)
            env_info = env.step(action)[brain_name]
            next_state = env_info.vector_observations[0]
            reward = env_info.rewards[0]
            done = env_info.local_done[0]
            agent.step(state, action, reward, next_state, done)
            state = next_state
            score += reward
            if done:
                break
        scores_win.append(score)
        scores.append(score)
        eps = max(eps_end, eps_decay * eps)
        print('\rEpisode {} \tAverage Score: {:.2f}'.format(e, np.mean(scores_win)), end='')
        if e % 100 == 0:
            print('\rEpisode {} \tAverage Score: {:.2f}'.format(e, np.mean(scores_win)))
        if np.mean(scores_win) >= 13.0:
            print('\nEnvironment solved in {:d} episodes! \tAverage Score: {:.2f}'.format(
                episodes, np.mean(scores_win)))
            torch.save(agent.qnet_local.state_dict(), 'checkpoint_navigation.pth')
            break
    return scores


In [14]: scores = dqn(episodes=episodes, horizon_t=horizon_t)

Episode 100      Average Score: 0.32
Episode 200      Average Score: 2.55
Episode 300      Average Score: 6.44
Episode 400      Average Score: 10.77
Episode 500      Average Score: 12.10
Episode 555      Average Score: 13.09
Environment solved in 456 episodes!      Average Score: 13.09

In [15]: # plot the scores
fig = plt.figure(figsize=(8,8))
ax = fig.add_subplot(111)
plt.plot(np.arange(len(scores)), scores)

```

```
plt.plot(pd.Series(scores).rolling(10).mean())  
plt.title('Score (Rewards)')  
plt.ylabel('Score')  
plt.xlabel('Episode #')  
plt.show()
```



```
In [16]: #when done, close the env  
env.close()
```

```
In [17]: ##
```